

this is basically <https://www.kaggle.com/bmarcos/image-recognition-gender-detection-inceptionv3> i only changed a couple things around so it would be able to run on my computer. the link has a lot more to look at to help understand the data, but i took it out to simplify what we actually need. i think it would be relatively simple to add in variables we want to test for, i just havent tried that. so far, this one is just training on the "male" variable.

i did have some trouble with the validation testing, so that still needs some work because when I ran it, the loss kept going up and im not sure why.

```
In [1]: import pandas as pd
import numpy as np
import cv2
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import f1_score
from keras import applications
from keras.applications.inception_v3 import InceptionV3, preprocess_input
from keras import optimizers
from keras.models import Sequential, Model
from keras.layers import Dropout, Flatten, Dense, GlobalAveragePooling2D
from keras.callbacks import ModelCheckpoint
from keras.preprocessing.image import ImageDataGenerator, array_to_img, img_to_array, load_img
from keras.utils import np_utils
from keras.optimizers import SGD

from IPython.core.display import display, HTML
from PIL import Image
from io import BytesIO
import base64

plt.style.use('ggplot')

%matplotlib inline
```

Using TensorFlow backend.

```
In [2]: import tensorflow as tf
print(tf.__version__)
```

2.1.0

## Hyperparameters

```
In [3]: # set variables
main_folder = 'Data/'
images_folder = main_folder + 'img_align_celeba/img_align_celeba/'

#EXAMPLE_PIC = images_folder + '000506.jpg'

TRAINING_SAMPLES = 10000
VALIDATION_SAMPLES = 2000
TEST_SAMPLES = 2000
IMG_WIDTH = 178
IMG_HEIGHT = 218
BATCH_SIZE = 16
NUM_EPOCHS = 15
```

```
In [4]: # import the data set that include the attribute for each picture
df_attr = pd.read_csv('Data/list_attr_celeba.csv')
df_attr.set_index('image_id', inplace=True)
df_attr.replace(to_replace=-1, value=0, inplace=True) #replace -1 by 0
df_attr.shape
#returns (202559, 40)
```

Out[4]: (202599, 40)

```
In [5]: #List of available attributes

# for i, j in enumerate(df_attr.columns):
#     print(i, j)
```

## Partitioning of data

```
In [6]: df_partition = pd.read_csv('Data/list_eval_partition.csv')
# display counter by partition
# 0 -> TRAINING
# 1 -> VALIDATION
# 2 -> TEST
df_partition['partition'].value_counts().sort_index()
```

Out[6]: 0      162770

```
1    19867
2    19962
Name: partition, dtype: int64
```

In [7]: *# join the partition with the attributes for GENDER*

```
df_partition.set_index('image_id', inplace=True)
df_par_attr = df_partition.join(df_attr['Male'], how='inner')
```

In [8]:

```
def load_reshape_img(fname):
    img = load_img(fname)
    x = img_to_array(img)/255.
    x = x.reshape((1,) + x.shape)

    return x

def generate_df(partition, attr, num_samples):
    '''
    partition
        0 -> train
        1 -> validation
        2 -> test
    '''

    df_ = df_par_attr[(df_par_attr['partition'] == partition)
                      & (df_par_attr[attr] == 0)].sample(int(num_samples/2))
    df_ = pd.concat([df_,
                     df_par_attr[(df_par_attr['partition'] == partition)
                                 & (df_par_attr[attr] == 1)].sample(int(num_samples/2))])

    # for Train and Validation
    if partition != 2:
        x_ = np.array([load_reshape_img(images_folder + fname) for fname in df_.index])
        x_ = x_.reshape(x_.shape[0], 218, 178, 3)
        y_ = np_utils.to_categorical(df_[attr],2)
    # for Test
    else:
        x_ = []
```

```

y_ = []

for index, target in df.iterrows():
    im = cv2.imread(images_folder + index)
    im = cv2.resize(cv2.cvtColor(im, cv2.COLOR_BGR2RGB), (IMG_WIDTH, IMG_HEIGHT)).astype(np.float)
    im = np.expand_dims(im, axis = 0)
    x_.append(im)
    y_.append(target[attr])

return x_, y_

```

### Pre-processing/data augmentation

```

In [9]: # Generate image generator for data augmentation
datagen = ImageDataGenerator(
    #preprocessing_function=preprocess_input,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True
)

# # Load one image and reshape
# img = load_img(EXAMPLE_PIC)
# x = img_to_array(img)/255.
# x = x.reshape((1,) + x.shape)

# # plot 10 augmented images of the loaded image
# plt.figure(figsize=(20,10))
# plt.suptitle('Data Augmentation', fontsize=28)

# i = 0
# for batch in datagen.flow(x, batch_size=1):
#     plt.subplot(3, 5, i+1)
#     plt.grid(False)
#     plt.imshow( batch.reshape(218, 178, 3))

#     if i == 9:
#         break

```

```
#     i += 1

# plt.show()
```

build data generators

```
In [10... # Train data
x_train, y_train = generate_df(0, 'Male', TRAINING_SAMPLES)

# Train - Data Preparation - Data Augmentation with generators
train_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
)

train_datagen.fit(x_train)

train_generator = train_datagen.flow(
    x_train, y_train,
    batch_size=BATCH_SIZE,
)
```

```
In [11... # Validation Data
x_valid, y_valid = generate_df(1, 'Male', VALIDATION_SAMPLES)

# Validation - Data Preparation - Data Augmentation with generators
valid_datagen = ImageDataGenerator(
    preprocessing_function=preprocess_input,
)

valid_datagen.fit(x_valid)

validation_generator = valid_datagen.flow(
```

```
x_valid, y_valid,
)
```

build model - for gender recognition

```
In [12... # Import InceptionV3 Model
inc_model = InceptionV3(weights='imagenet',
                        include_top=False,
                        input_shape=(IMG_HEIGHT, IMG_WIDTH, 3))

print("number of layers:", len(inc_model.layers))
#inc_model.summary()
```

Downloading data from [https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception\\_v3\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://github.com/fchollet/deep-learning-models/releases/download/v0.5/inception_v3_weights_tf_dim_ordering_tf_kernels_notop.h5)  
87916544/87910968 [=====] - 2s 0us/step  
number of layers: 311

```
In [13... #Adding custom Layers
x = inc_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation = "relu")(x)
x = Dropout(0.5)(x)
x = Dense(512, activation = "relu")(x)
predictions = Dense(2, activation = "softmax")(x)
```

```
In [14... # creating the final model
model_ = Model(inputs=inc_model.input, outputs=predictions)

# Lock initial layers to do not be trained
for layer in model_.layers[:52]:
    layer.trainable = False

# compile the model
model_.compile(optimizer=SGD(lr=0.0001, momentum=0.9)
               , loss='categorical_crossentropy'
               , metrics=['accuracy'])
```

train model

```
In [15... #https://keras.io/models/sequential/ fit generator
```

```
checkpointer = ModelCheckpoint(filepath='weights.best.inc.male.hdf5',
                               verbose=1, save_best_only=True)
```

```
In [16... hist = model_.fit_generator(train_generator
                                , validation_data = (x_valid, y_valid)
                                , steps_per_epoch= TRAINING_SAMPLES/BATCH_SIZE
                                , epochs= NUM_EPOCHS
                                , callbacks=[checkpointer]
                                , verbose=1
                                )
```

Epoch 1/15

625/625 [=====] - 694s 1s/step - loss: 0.5346 - accuracy: 0.7277 - val\_loss: 0.3153 - val\_accuracy: 0.8680

Epoch 00001: val\_loss improved from inf to 0.31532, saving model to weights.best.inc.male.hdf5

Epoch 2/15

625/625 [=====] - 691s 1s/step - loss: 0.3269 - accuracy: 0.8657 - val\_loss: 0.2636 - val\_accuracy: 0.8825

Epoch 00002: val\_loss improved from 0.31532 to 0.26362, saving model to weights.best.inc.male.hdf5

Epoch 3/15

625/625 [=====] - 688s 1s/step - loss: 0.2710 - accuracy: 0.8865 - val\_loss: 0.2285 - val\_accuracy: 0.9010

Epoch 00003: val\_loss improved from 0.26362 to 0.22854, saving model to weights.best.inc.male.hdf5

Epoch 4/15

625/625 [=====] - 679s 1s/step - loss: 0.2363 - accuracy: 0.9038 - val\_loss: 0.1991 - val\_accuracy: 0.9160

Epoch 00004: val\_loss improved from 0.22854 to 0.19912, saving model to weights.best.inc.male.hdf5

Epoch 5/15

625/625 [=====] - 683s 1s/step - loss: 0.2210 - accuracy: 0.9081 - val\_loss: 0.1964 - val\_accuracy: 0.9210

Epoch 00005: val\_loss improved from 0.19912 to 0.19639, saving model to weights.best.inc.male.hdf5

Epoch 6/15

625/625 [=====] - 668s 1s/step - loss: 0.2005 - accuracy: 0.9227 - val\_loss: 0.1941 - val\_accuracy: 0.9190

Epoch 00006: val\_loss improved from 0.19639 to 0.19409, saving model to weights.best.inc.male.hdf5

Epoch 7/15

625/625 [=====] - 670s 1s/step - loss: 0.1874 - accuracy: 0.9256 - val\_loss: 0.231 - val\_accuracy: 0.9100

Epoch 00007: val\_loss did not improve from 0.19409

Epoch 8/15

625/625 [=====] - 678s 1s/step - loss: 0.1846 - accuracy: 0.9254 - val\_loss: 0.1997 - val\_accuracy: 0.9230

Epoch 00008: val\_loss did not improve from 0.19409

Epoch 9/15

625/625 [=====] - 685s 1s/step - loss: 0.1600 - accuracy: 0.9353 - val\_loss: 0.1771 - val\_accuracy: 0.9315

Epoch 00009: val\_loss improved from 0.19409 to 0.17715, saving model to weights.best.inc.male.hdf5

Epoch 10/15

625/625 [=====] - 672s 1s/step - loss: 0.1550 - accuracy: 0.9396 - val\_loss: 0.2427 - val\_accuracy: 0.9015

Epoch 00010: val\_loss did not improve from 0.17715

Epoch 11/15

625/625 [=====] - 681s 1s/step - loss: 0.1478 - accuracy: 0.9426 - val\_loss: 0.1952 - val\_accuracy: 0.9225

Epoch 00011: val\_loss did not improve from 0.17715

Epoch 12/15

625/625 [=====] - 695s 1s/step - loss: 0.1464 - accuracy: 0.9437 - val\_loss: 0.2124 - val\_accuracy: 0.9145

Epoch 00012: val\_loss did not improve from 0.17715

Epoch 13/15

625/625 [=====] - 704s 1s/step - loss: 0.1330 - accuracy: 0.9486 - val\_loss: 0.1887 - val\_accuracy: 0.9260

Epoch 00013: val\_loss did not improve from 0.17715

Epoch 14/15

625/625 [=====] - 713s 1s/step - loss: 0.1250 - accuracy: 0.9522 - val\_loss: 0.1573 - val\_accuracy: 0.9380

Epoch 00014: val\_loss improved from 0.17715 to 0.15735, saving model to weights.best.inc.male.hdf5


Epoch 15/15




625/625 [=====] - 666s 1s/step - loss: 0.1172 - accuracy: 0.9536 - val\_loss: 0.1922 - val\_accuracy: 0.9245

Epoch 00015: val\_loss did not improve from 0.15735

```
In [17... # Plot loss function value through epochs
plt.figure(figsize=(18, 4))
plt.plot(hist.history['loss'], label = 'train')
plt.plot(hist.history['val_loss'], label = 'valid')
plt.legend()
plt.title('Loss Function')
plt.show()
```



```
In [18... # Plot accuracy through epochs
plt.figure(figsize=(18, 4))
plt.plot(hist.history['accuracy'], label = 'train')
plt.plot(hist.history['val_accuracy'], label = 'valid')
plt.legend()
plt.title('Accuracy')
plt.show()
```



model evaluation

```
In [19... #load the best model
model_.load_weights('weights.best.inc.male.hdf5')
```

```
In [20... # Test Data
x_test, y_test = generate_df(2, 'Male', TEST_SAMPLES)

# generate prediction
model_predictions = [np.argmax(model_.predict(feature)) for feature in x_test ]

# report test accuracy
test_accuracy = 100 * np.sum(np.array(model_predictions)==y_test) / len(model_predictions)
print('Model Evaluation')
print('Test accuracy: %.4f%%' % test_accuracy)
print('f1_score:', f1_score(y_test, model_predictions))
```

Model Evaluation

Test accuracy: 93.3500%

f1\_score: 0.9290666666666666

```
In [21... #dictionary to name the prediction
gender_target = {0: 'Female'
                 , 1: 'Male'}

def img_to_display(filename):
```

```

i = Image.open(filename)
i.thumbnail((200, 200), Image.LANCZOS)

with BytesIO() as buffer:
    i.save(buffer, 'jpeg')
    return base64.b64encode(buffer.getvalue()).decode()

##this part is extra####
def display_result(filename, prediction, target):
    '''
    Display the results in HTML

    '''

    gender = 'Male'
    gender_icon = "https://i.imgur.com/nxWan2u.png"

    if prediction[1] <= 0.5:
        gender_icon = "https://i.imgur.com/oAAb8rd.png"
        gender = 'Female'

    display_html = '''
    <div style="overflow: auto; border: 2px solid #D8D8D8;
        padding: 5px; width: 420px;" >
        
        <div style="padding: 10px 0px 0px 20px; overflow: auto;">
            
            <h3 style="margin-left: 50px; margin-top: 2px;">{}/h3>
            <p style="margin-left: 50px; margin-top: -6px; font-size: 12px">{} prob.</p>

        </div>
    </div>
    '''.format(img_to_display(filename)
        , gender_icon
        , gender
        , "{0:.2f}%".format(round(max(prediction)*100,2))
        , gender_target[target]
        , filename.split('/')[-1]
        )

    display(HTML(display_html))

```

```
In [22... def gender_prediction(filename):  
    '''  
    predict the gender  
  
    input:  
        filename: str of the file name  
  
    return:  
        array of the prob of the targets.  
  
    '''  
  
    im = cv2.imread(filename)  
    im = cv2.resize(cv2.cvtColor(im, cv2.COLOR_BGR2RGB), (178, 218)).astype(np.float32) / 255.0  
    im = np.expand_dims(im, axis = 0)  
  
    # prediction  
    result = model_.predict(im)  
    prediction = np.argmax(result)  
  
    return result
```

```
In [23... #select random images of the test partition  
df_to_test = df_par_attr[(df_par_attr['partition'] == 2)].sample(2)  
  
for index, target in df_to_test.iterrows():  
    result = gender_prediction(images_folder + index)  
  
    #display result  
    display_result(images_folder + index, result[0], target['Male'])
```



**Female**

100.00% prob.



**Female**

99.99% prob.